

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, models, transforms
import os
from tqdm import tqdm # Import tqdm for progress bars
import copy
from sklearn.metrics import accuracy_score

# Set device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

# 1. Load Pre-trained Model and Modify Output Layer
def load_and_modify_model(num_classes):
    """Loads a pre-trained model (resnet18), replaces the classifier,
    and returns the modified model."""
    model = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)
# Using ResNet18
    num_fts = model.fc.in_features
    model.fc = nn.Linear(num_fts, num_classes)
    return model.to(device)

# 2. Define Data Transformations
def get_data_transforms(augment=False):
    """Defines data transformations, with optional augmentation."""
    if augment:
        train_transforms = transforms.Compose([
            transforms.RandomResizedCrop(224),
            transforms.RandomHorizontalFlip(),
            transforms.RandomRotation(degrees=15),
            transforms.ColorJitter(brightness=0.2, contrast=0.2,
saturation=0.2),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
std=[0.229, 0.224, 0.225])
        ])
    else:
        train_transforms = transforms.Compose([
            transforms.Resize(256),
            transforms.CenterCrop(224),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
std=[0.229, 0.224, 0.225])
        ])

    test_transforms = transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),

```

```

        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225]))
    ])

    return train_transforms, test_transforms

# 3. Load Datasets and Create DataLoaders
def load_datasets(data_dir, train_transforms, test_transforms,
batch_size):
    """Loads datasets, creating data loaders."""
    train_dataset = datasets.ImageFolder(os.path.join(data_dir,
'train'), transform=train_transforms)
    test_dataset = datasets.ImageFolder(os.path.join(data_dir,
'test'), transform=test_transforms)

    train_loader = DataLoader(train_dataset, batch_size=batch_size,
shuffle=True, num_workers=4)
    test_loader = DataLoader(test_dataset, batch_size=batch_size,
shuffle=False, num_workers=4)

    return train_loader, test_loader, len(train_dataset.classes)

# 4. Define Training and Evaluation Functions
def train_model(model, train_loader, criterion, optimizer, num_epochs,
augment=False):
    """Trains the model, printing loss and accuracy, and returns best
performing model."""
    best_model_wts = copy.deepcopy(model.state_dict())
    best_accuracy = 0.0

    for epoch in range(num_epochs):
        print(f'Epoch {epoch+1}/{num_epochs}')
        print('-' * 10)

        model.train() # Set model to training mode

        running_loss = 0.0
        all_labels = []
        all_predictions = []

        for inputs, labels in tqdm(train_loader, desc="Training",
unit="batch"):
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

```

```

        running_loss += loss.item() * inputs.size(0)
        _, predicted = torch.max(outputs, 1) # Get predictions
        all_labels.extend(labels.cpu().numpy())
        all_predictions.extend(predicted.cpu().numpy())

    epoch_loss = running_loss / len(train_loader.dataset)
    epoch_accuracy = accuracy_score(all_labels, all_predictions)

    print(f'Train Loss: {epoch_loss:.4f} Train Acc:
{epoch_accuracy:.4f}')

    # Evaluate and save the best model if accuracy is better
    accuracy = evaluate_model(model, test_loader)
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_model_wts = copy.deepcopy(model.state_dict())
        print(f'Validation accuracy improved
({best_accuracy:.4f}), saving model.')

    print(f'Best val accuracy: {best_accuracy:4f}')
    model.load_state_dict(best_model_wts)
    return model

def evaluate_model(model, test_loader):
    """Evaluates the model and returns the accuracy."""
    model.eval() # Set model to evaluation mode
    all_labels = []
    all_predictions = []

    with torch.no_grad():
        for inputs, labels in tqdm(test_loader, desc="Testing",
unit="batch"):
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, predicted = torch.max(outputs, 1) # Get predictions
            all_labels.extend(labels.cpu().numpy())
            all_predictions.extend(predicted.cpu().numpy())

    accuracy = accuracy_score(all_labels, all_predictions)
    print(f'Validation Accuracy: {accuracy:.4f}')
    return accuracy

# 5. Main Execution
if __name__ == '__main__':
    # Parameters
    data_dir = 'EuroSAT_RGB' # Path to the dataset
    num_epochs = 10 # Increased for more training
    batch_size = 32

```


[illegible]

Validation Accuracy: 0.9482
Best val accuracy: 0.948200